Postgres-XC

July 12th, 2011
Koichi Suzuki
NTT DATA INTELLILINK CORPORATION

Overview of Postgres-XC

Symmetric PostgreSQL cluster

- No Master
- No Slave
  - No READ ONLY slaves
  - Every node can issue both READ/WRITE
- Transparent Transaction Management

Now Version 0.9.5

- Generally available next calendar year

PostgreSQL Master/Slave with Log Shipping

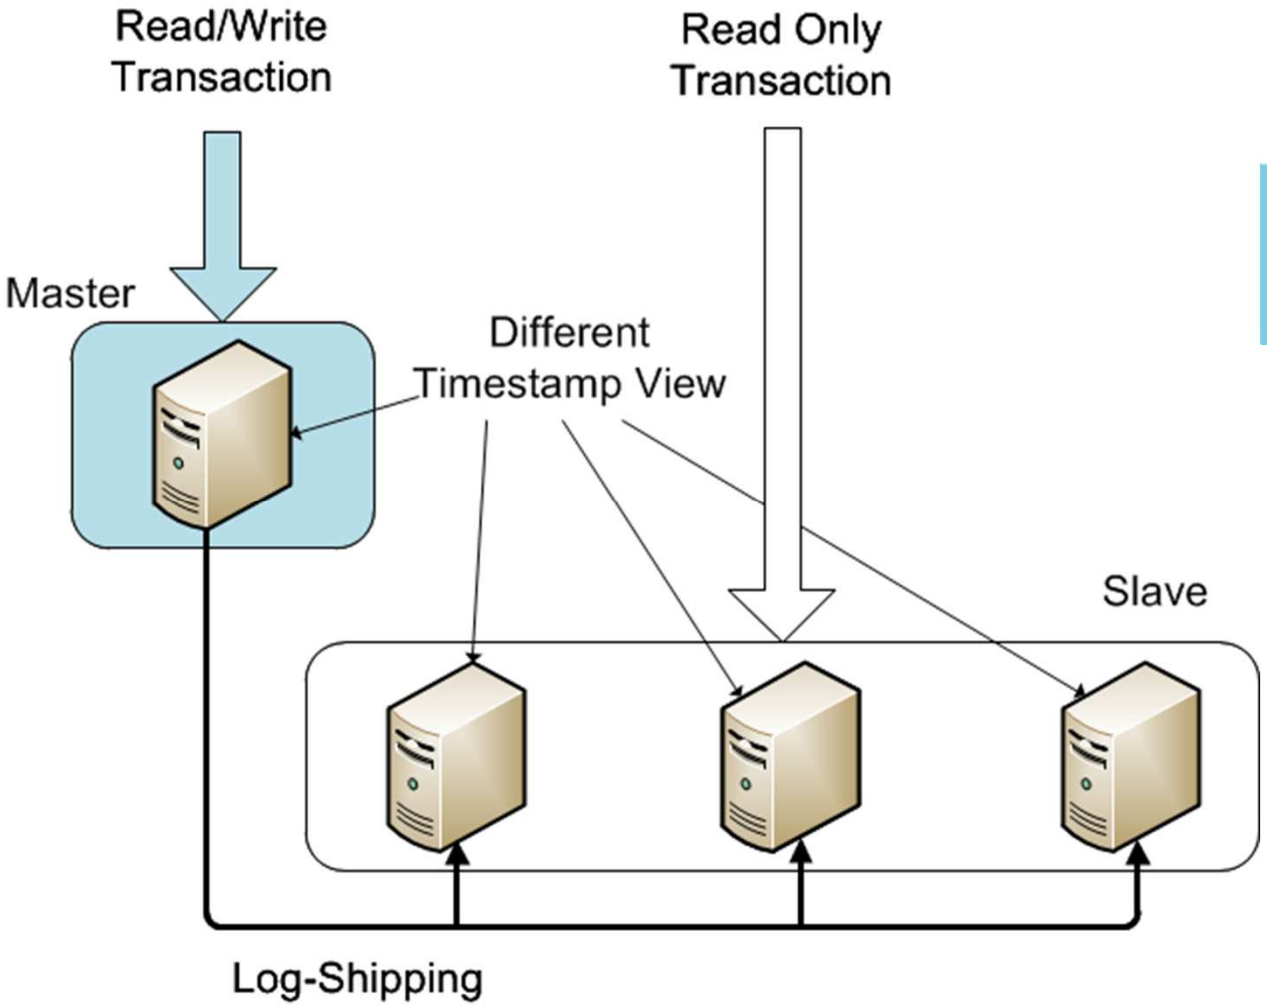

Postgres-XC Symmetric Cluster

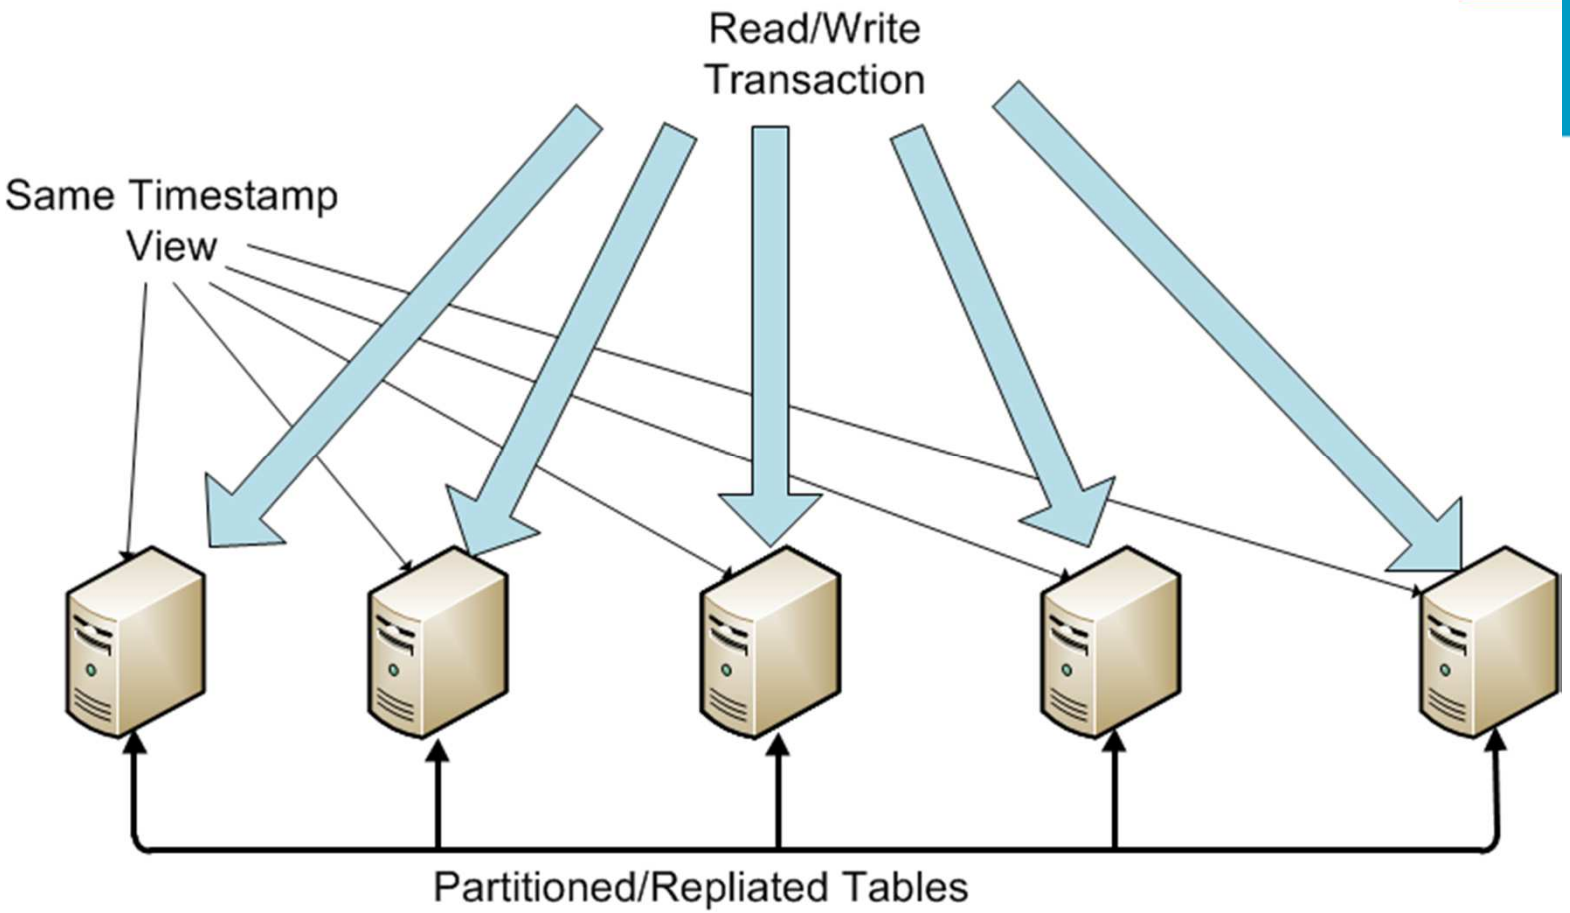

Server Configuration and GTM-Proxy

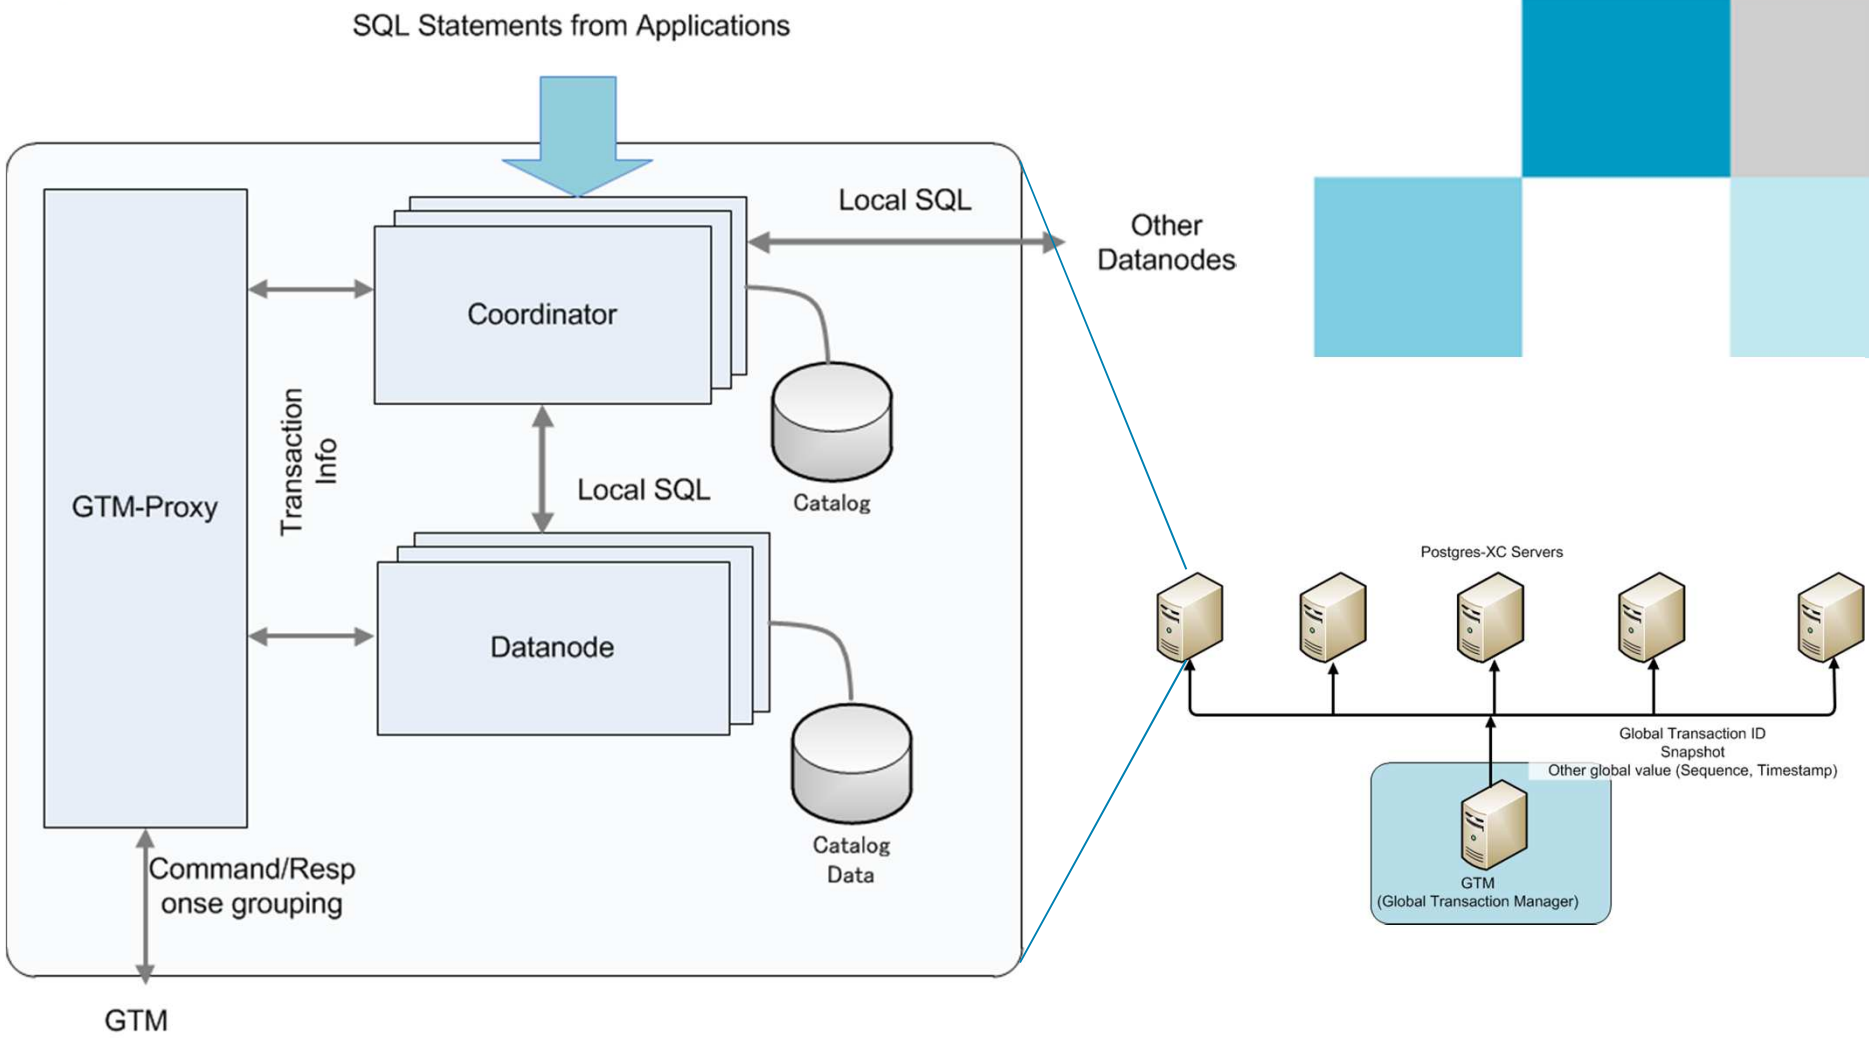

Scalability

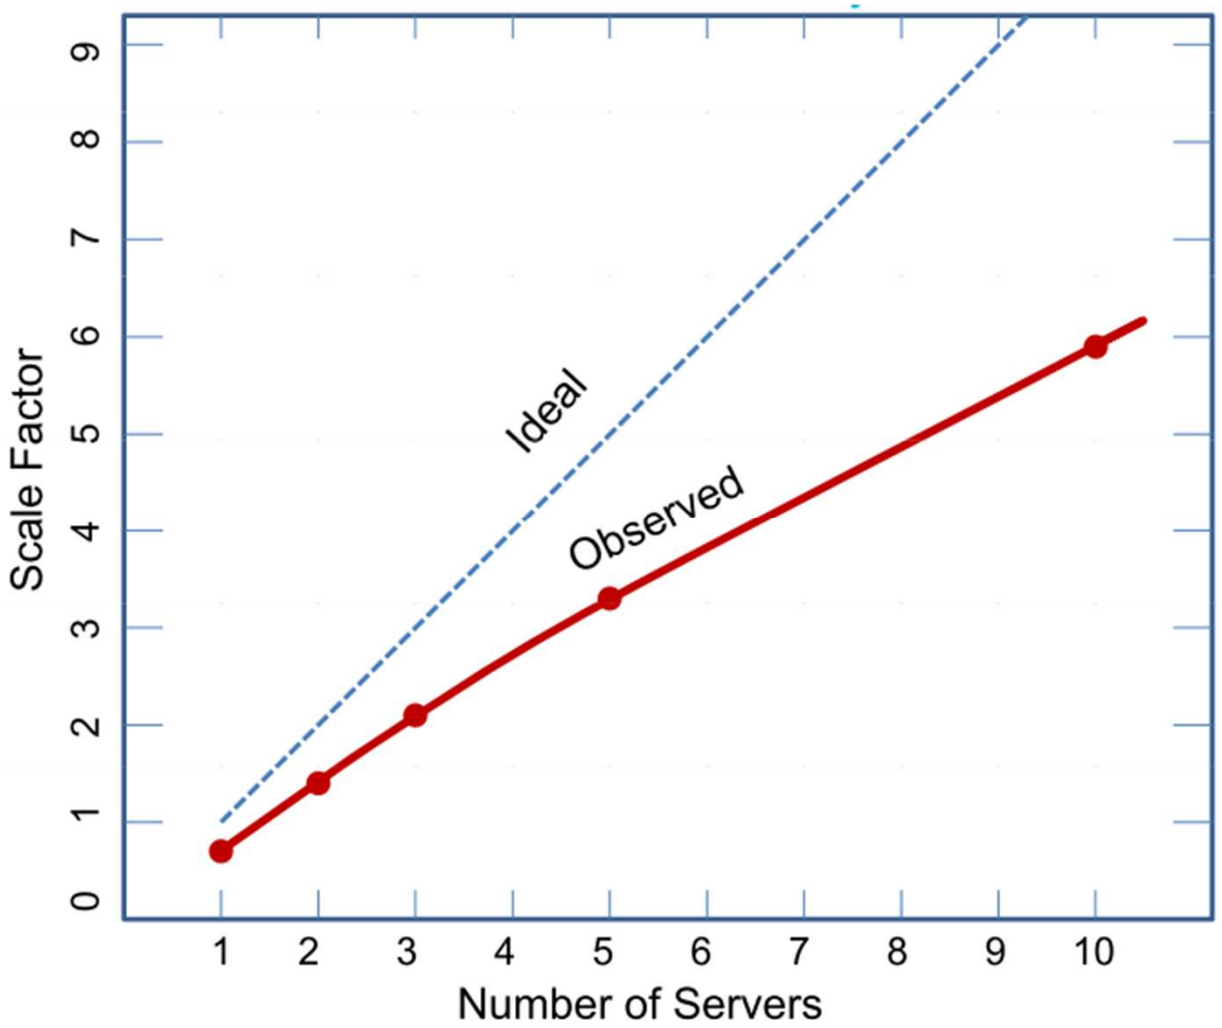

DBT-1 (Rev)

Current Status

- Now V 0.9.5 is available
- License changed to PostgreSQL license
  - Free to bring outcome back to PostgreSQL

GTM: Key for Transaction Transparency

- Consistent Transaction ID (GXID) throughout the system
- Provide global snapshot for consistent visibility from any server

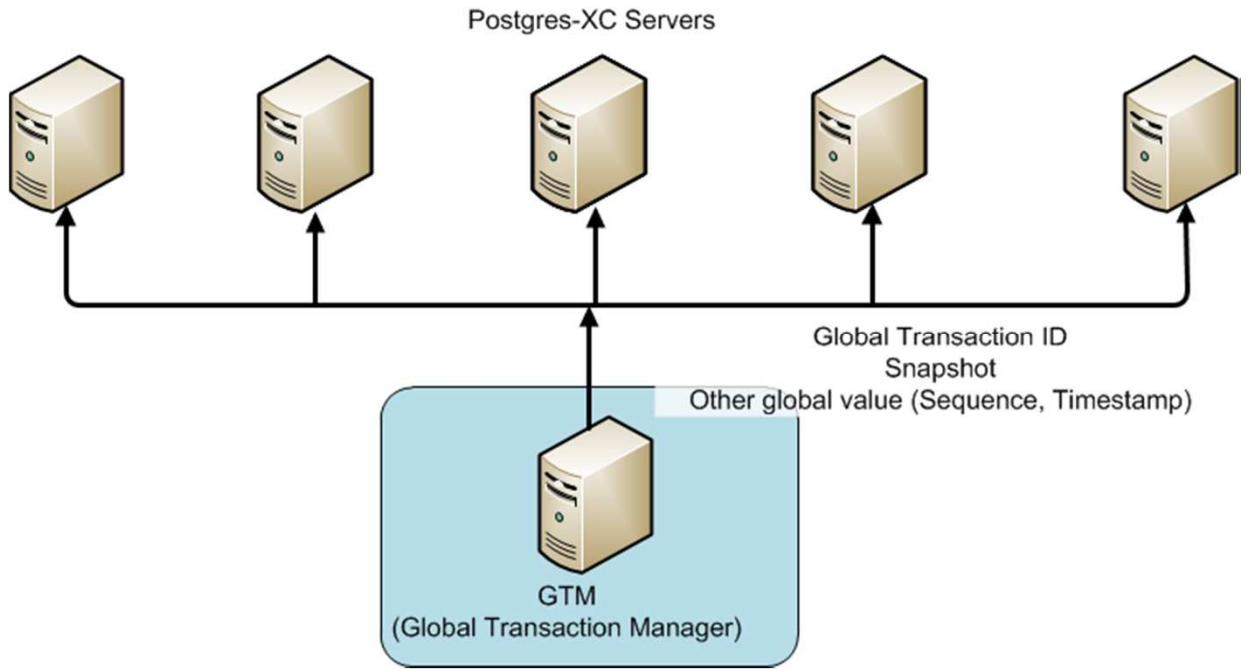

Requirements Since Last Year …

Solution for GTM as SPOF

- GTM Standby

Support same SQL statements as original PostgreSQL

- Functions
- Views
- Cross-node joins
- Role/User/Tablespace
- Transparent DDLs
- Many others

Other High Availability Feature such as

- Data Node Standby
- Consistent Backup and Recovery

Flexible Node Configuration

- On-line addition/Removal

GTM Standby

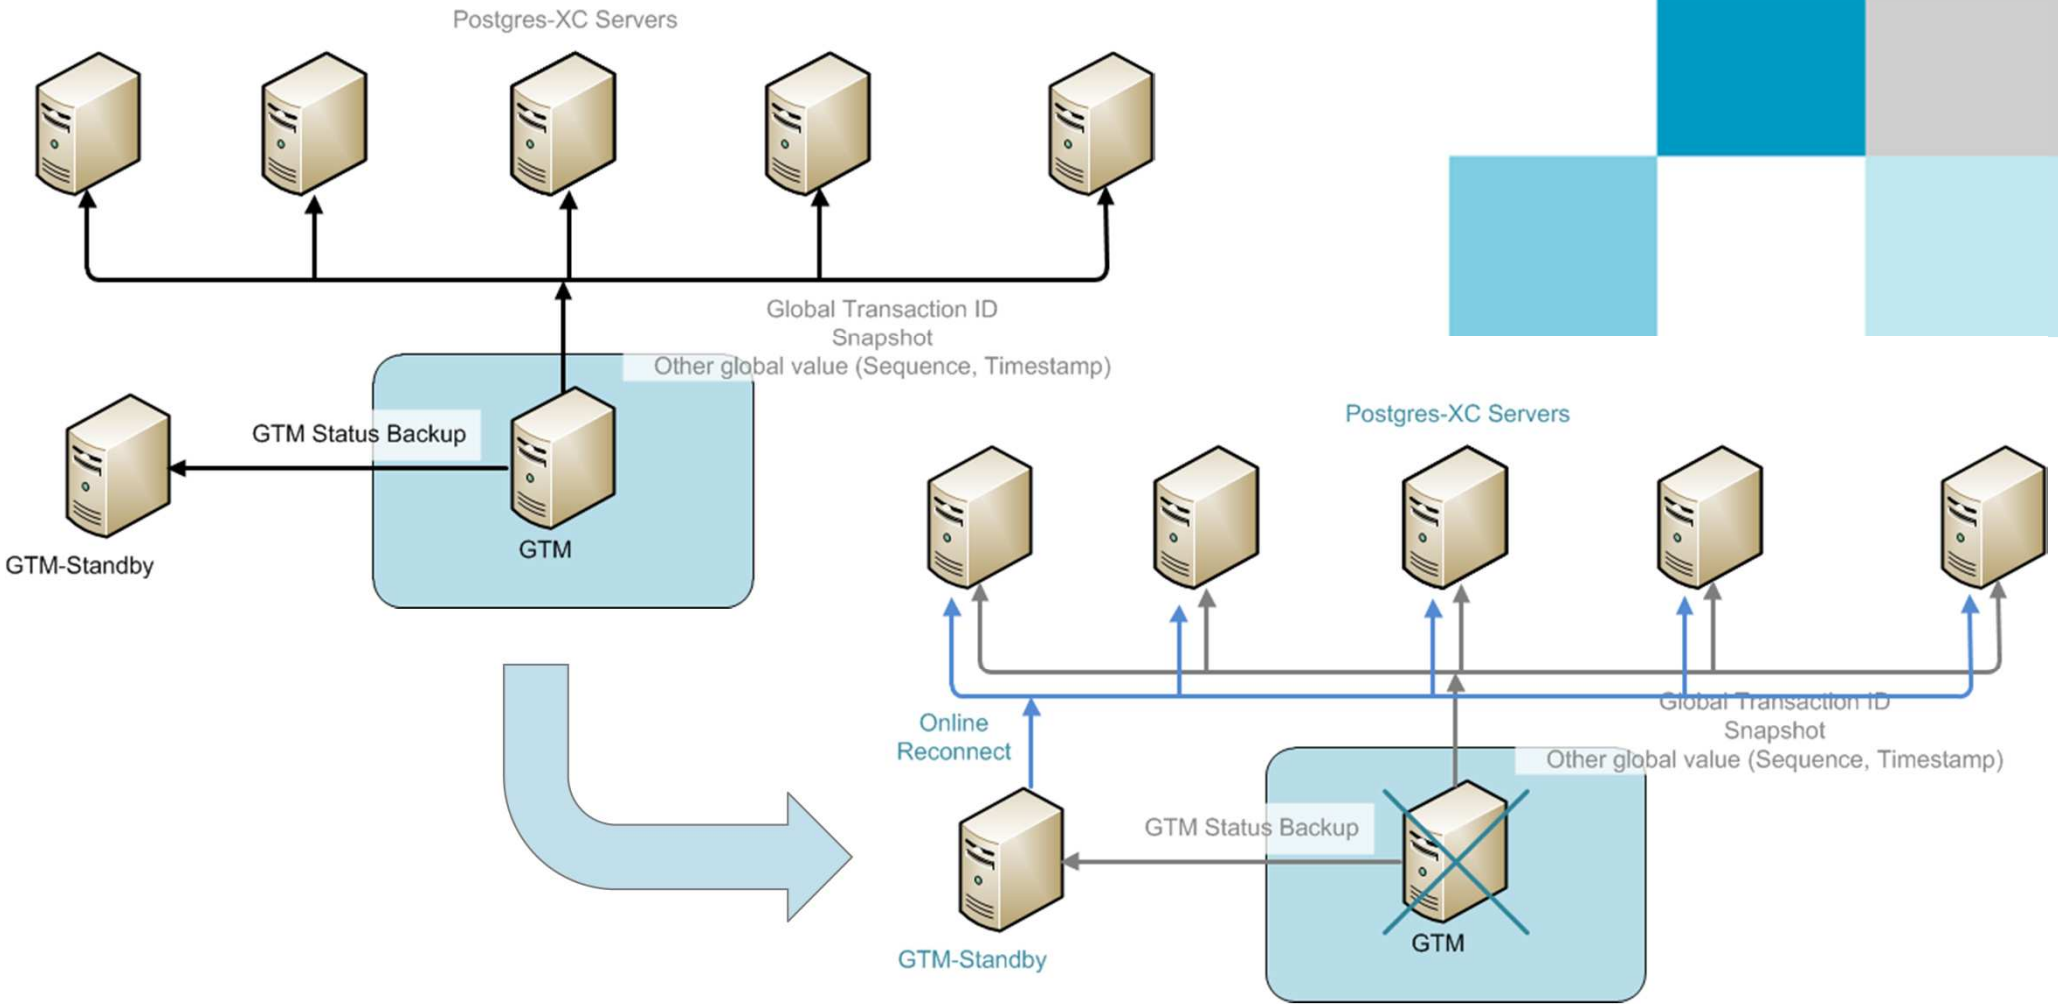

GTM Standby Requirements

Online Promote and Reconnect

- Invisible from applications
  - Can be visible from GTM-Proxy
- Transactions should be able to continue to run

GTM-Standby: Current Status

- Infrastructure Available: V 0.9.5
- Improvement in progress
  - Connect to GTM at anytime
    - At present, GTM-Standby should be the first to connect to GTM
  - Get rid of any chance of backup information loss
    - Backup first
    - Negotiate the last message at reconnect
  - Performance
    - Backup grouping and decrease response
- Improvement scheduled at the next release

Postgres-XC Statement Extension

Analyze incoming statements

Determine which datanodes are to be involved

- Compose Local Statement for each Datanode

Datanode handle the local statement.

- Result to the coordinator

If needed, materialize the result at the coordinator and calculate the final result.

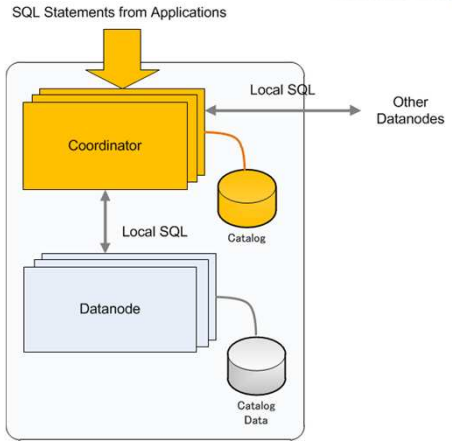

Postgres-XC Statement Extension

Analyze incoming statements

Determine which datanodes are to be involved

- Compose Local Statement for each Datanode

Datanode handle the local statement.

- Result to the coordinator

If needed, materialize the result at the coordinator and calculate the final result.

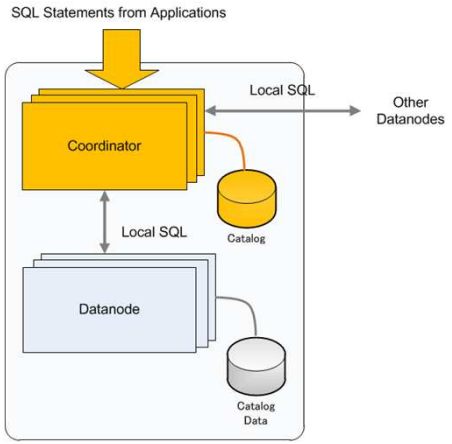

Postgres-XC Statement Extension

Analyze incoming statements

Determine which datanodes are to be involved

- Compose Local Statement for each Datanode

Datanode handle the local statement.

- Result to the coordinator

If needed, materialize the result at the coordinator and calculate the final result.

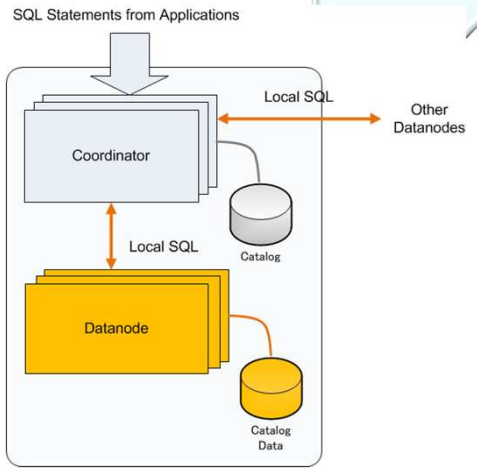

Postgres-XC Statement Extension

Analyze incoming statements

Determine which datanodes are to be involved

- Compose Local Statement for each Datanode

Datanode handle the local statement.

- Result to the coordinator

If needed, materialize the result at the coordinator and calculate the final result.

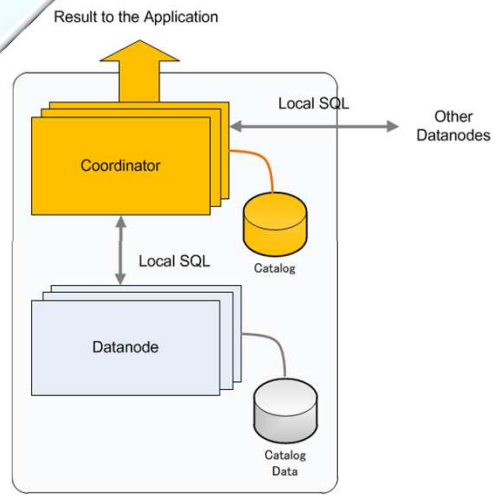

Optimizing Statements (V 0.9.5)

Push-down as many clause as possAZible

- Join
- WHERE Clause
- Aggregate
- Functions (when used in WHERE clause)
- Column projection

Uses the following information

- If each table is replicated or partitioned
- Partition key
- Partition algorism (Hash/Modulo/Round Robin)

Future Improvement

Candidate

- Use statistic info.
- Use Semi-Join to determine joining rows
- Direct join tuple transfer among datanodes
- Much more …

XC Optimization Examples (Join-1)

- Both Tables Are Replicated

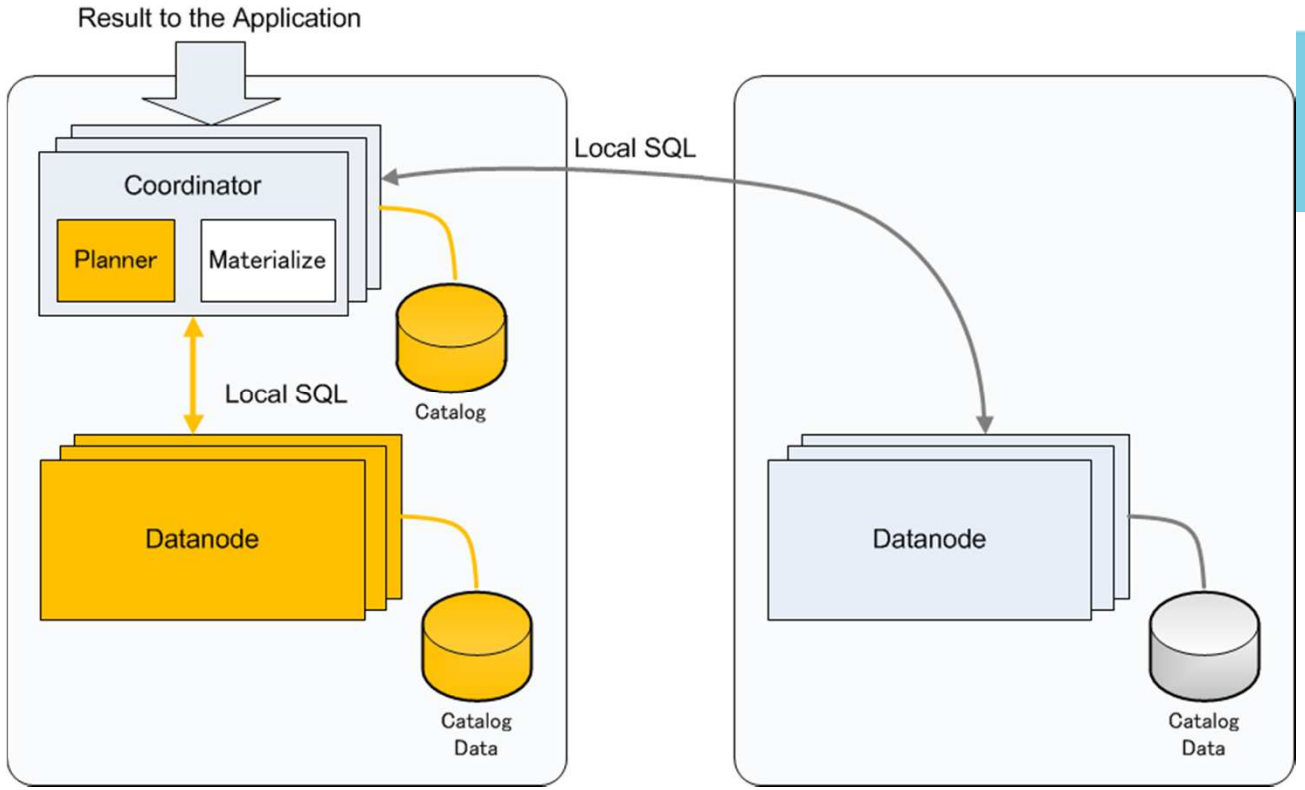

XC Optimization Examples (Join-2)

- Replicated Table and Partitioned Table
  - Cannot determine which datanode to go

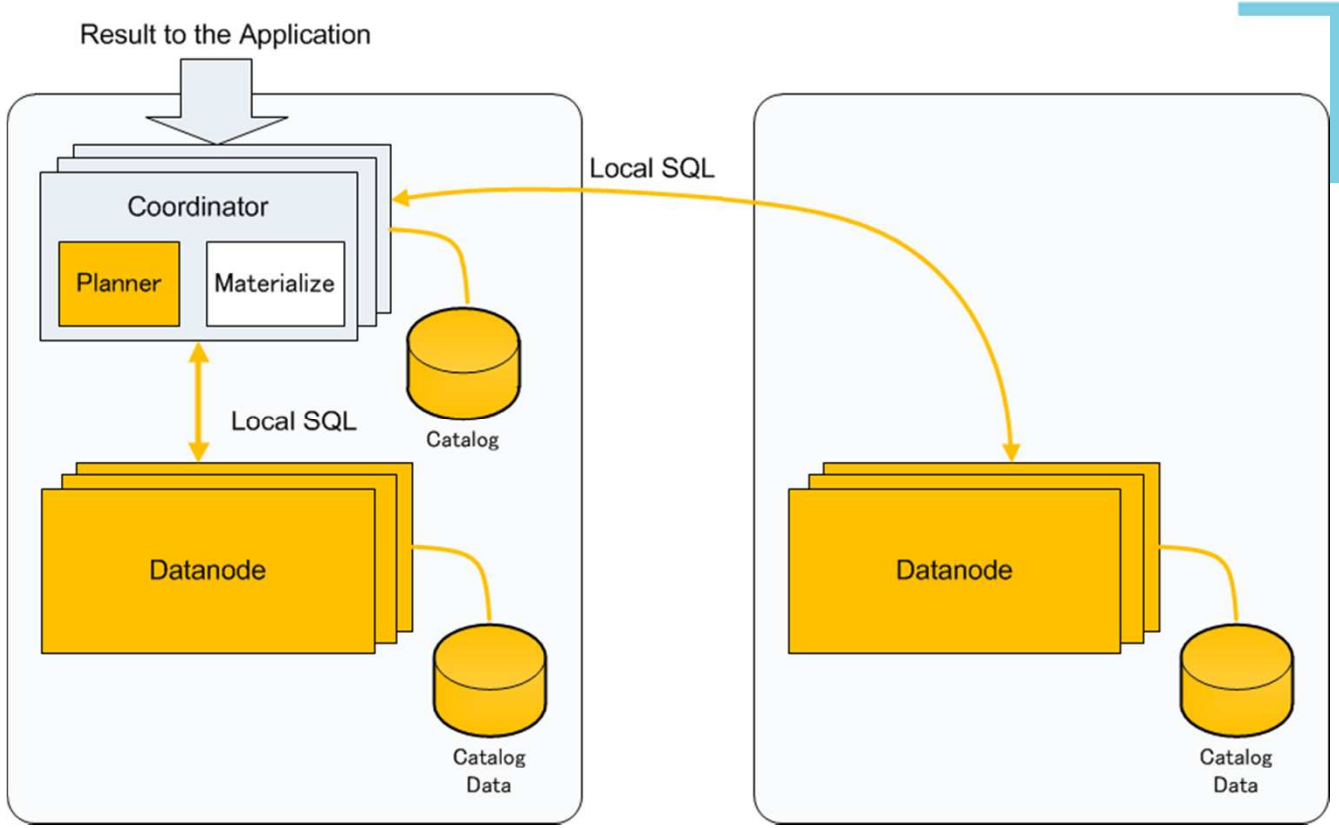

XC Optimization Examples (Join-3)

- Replicated Table and Partitioned Table
  - Can determine which datanode to go from WHERE clause

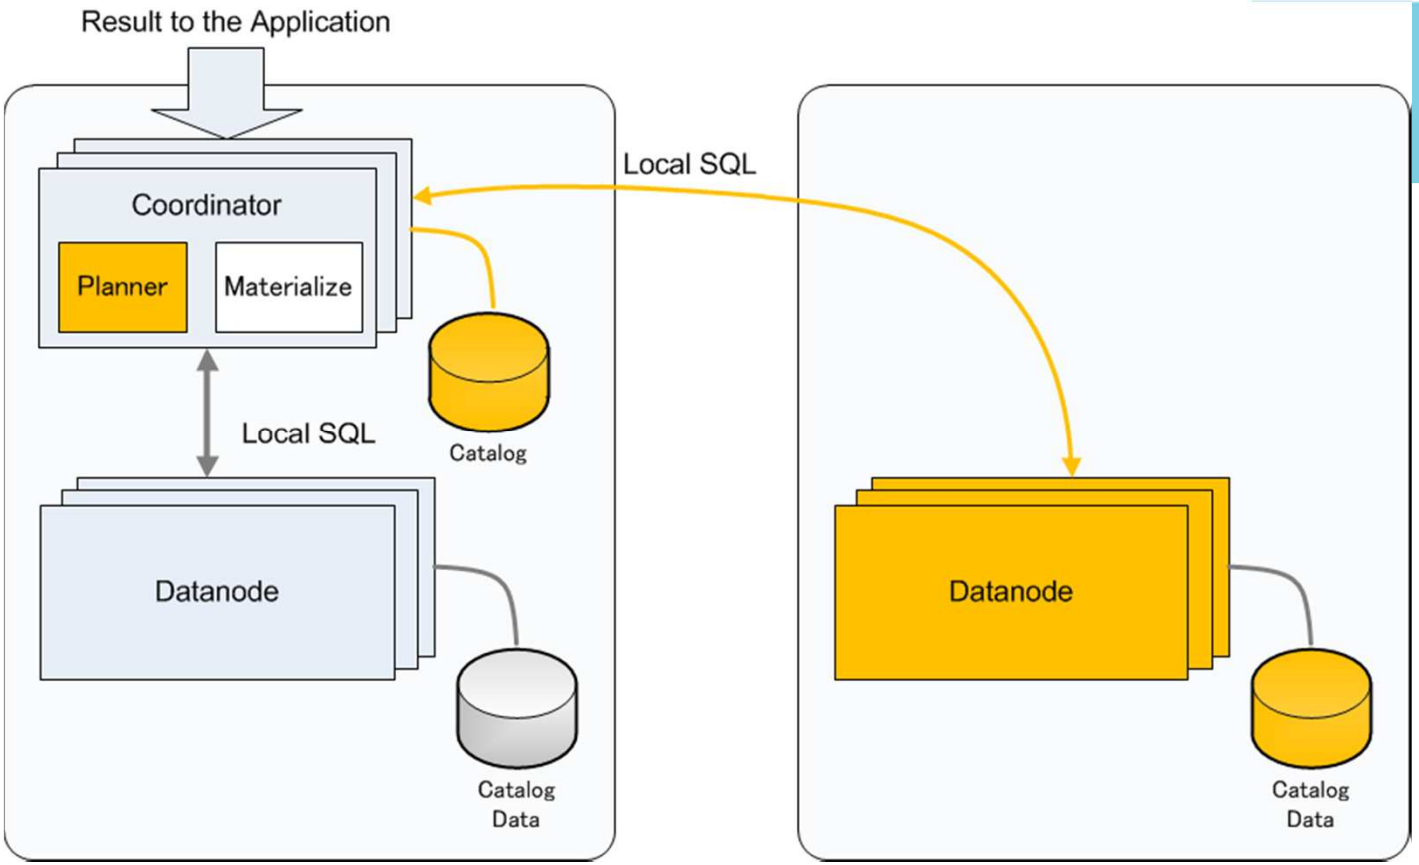

XC Optimization Examples (Join-4)

- Partitioned Table and Partitioned Table
  - Both Join columns are distribution (partitioning) column
  - Where clause can determine which datanode to go

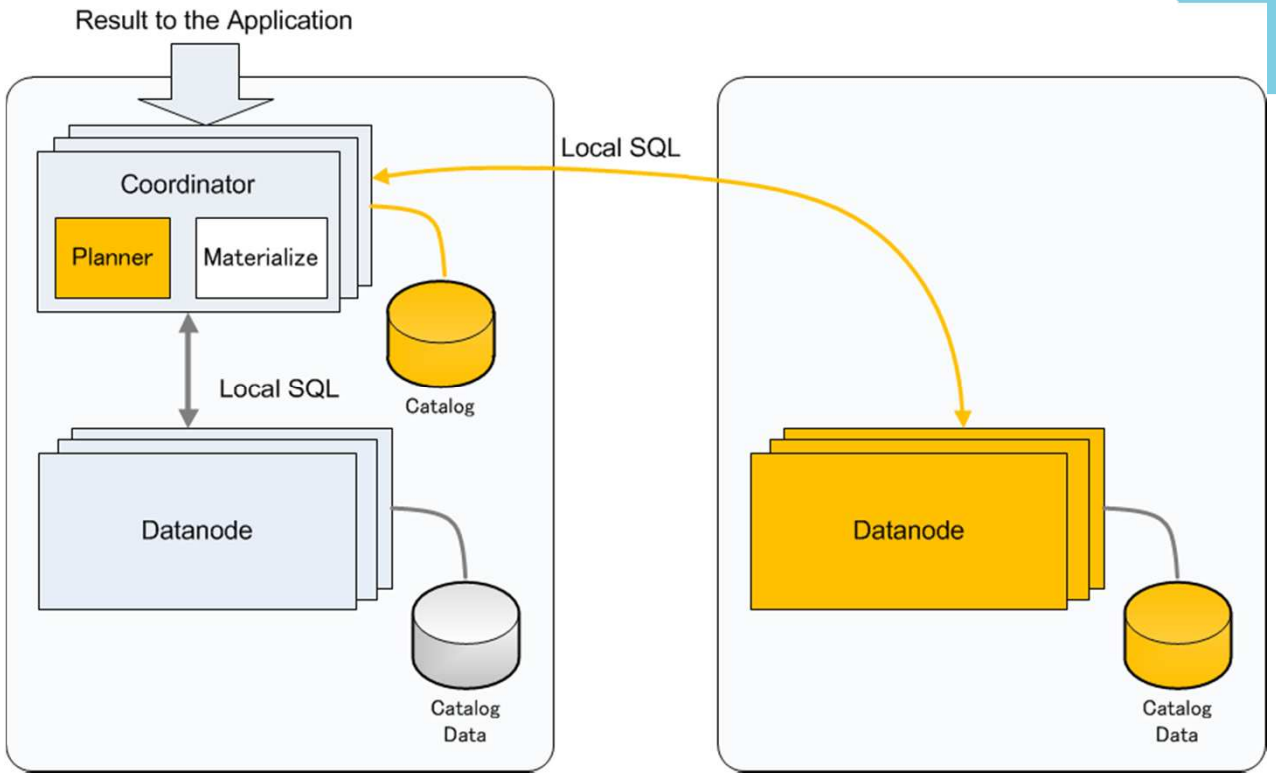

XC Optimization Examples (Join-5)

- Partitioned Table and Partitioned Table
  - Both Join columns are distribution (partitioning) column

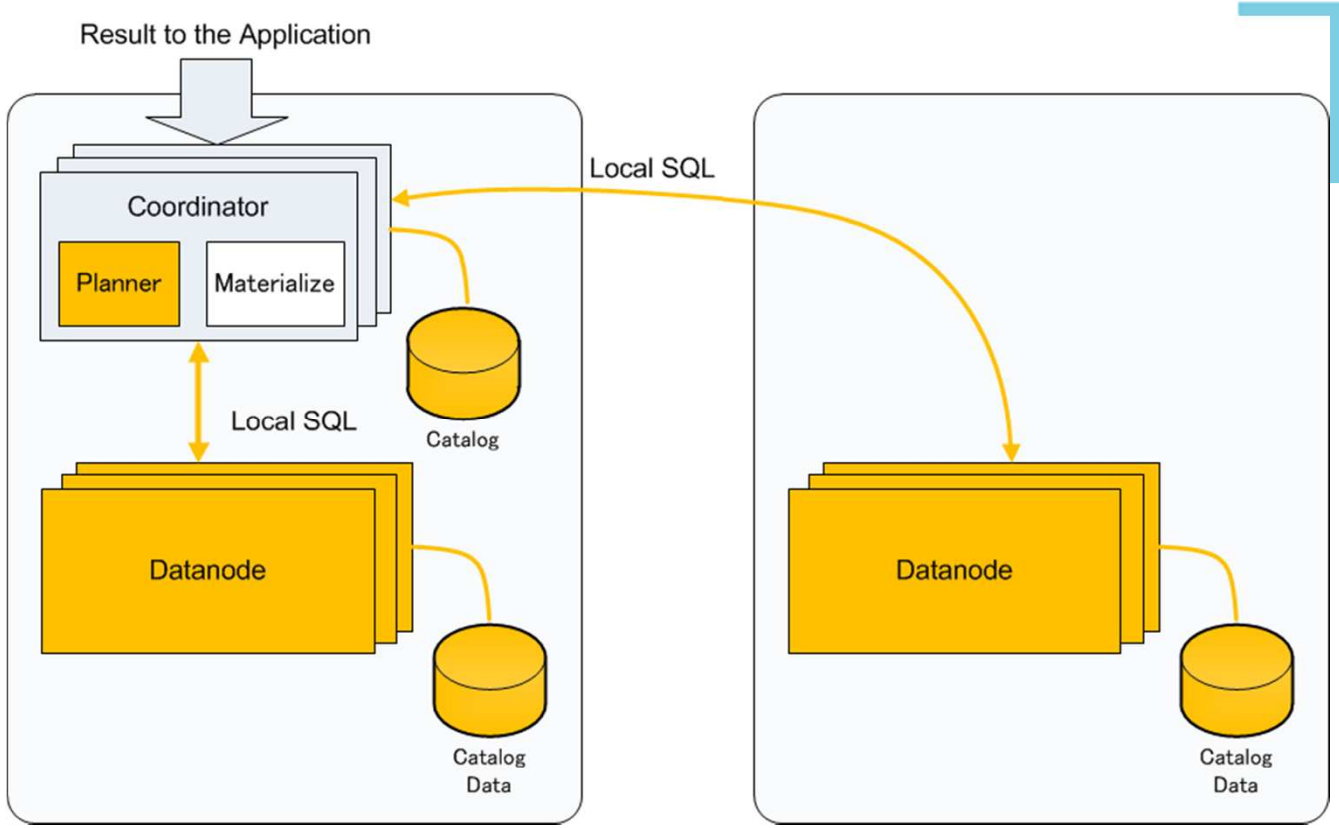

XC Optimization Examples (Join-6)

- Partitioned Table and Partitioned Table
  - One of Join columns are not distribution (partitioning) column

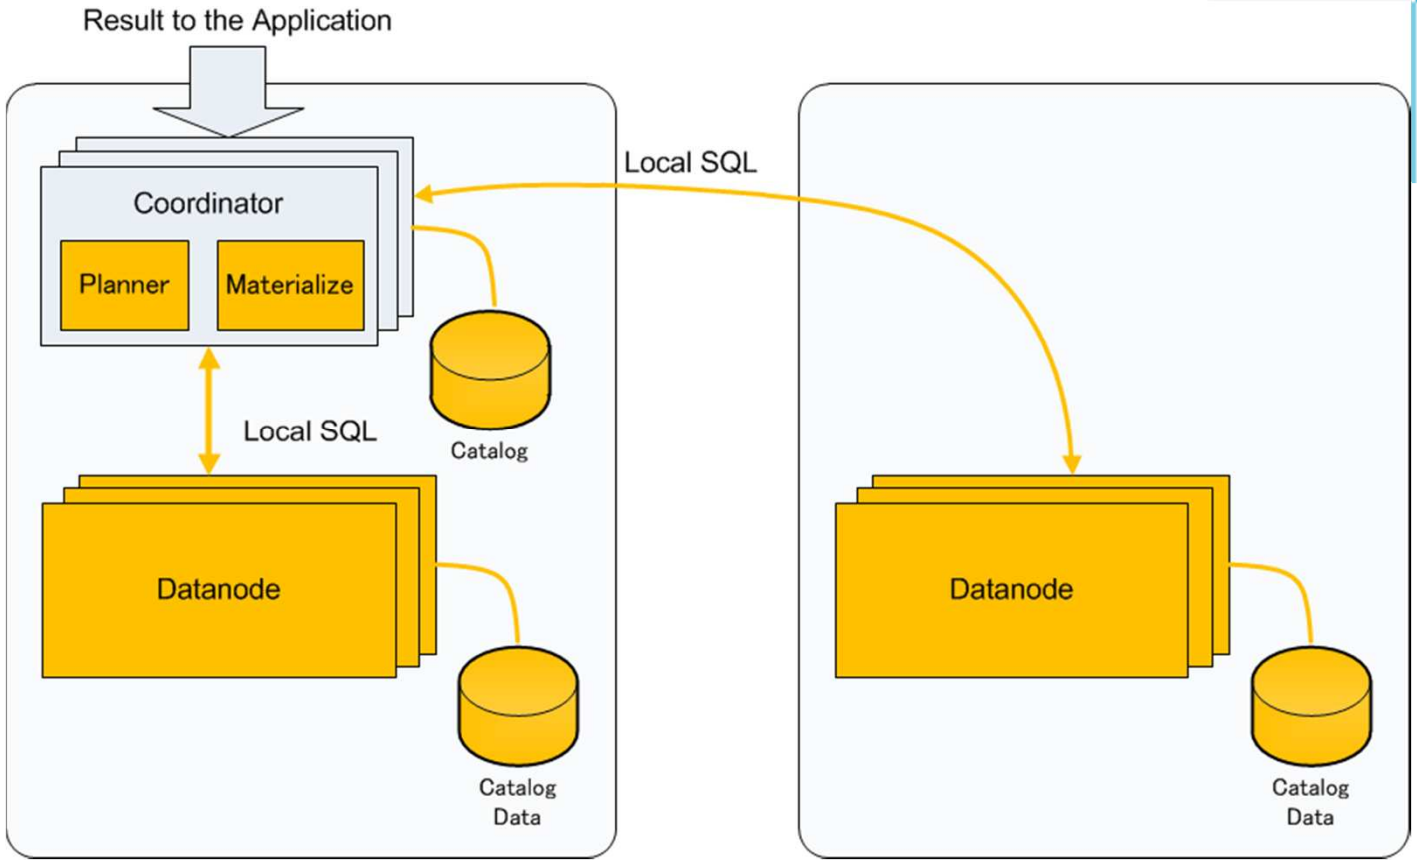

XC Statement Handling Summary

- Now can handle wide variety of PostgreSQL statement.
- Still in progress
  - HAVING
  - PREPARE, EXECUTE, CURSOR
    - Eliminate restrictions
  - WITH/WITH RECURSIVE
  - General Subqueries
  - Functions with more than one statement
  - SELECT INTO (CREATE TABLE AS)
  - Triggers
  - Temp tables
- Challenges
  - Global constraint
  - More Optimization
  - More Parallelism
- Miscellaneous
  - LISTEN/NOTIFY/UNLISTEN

Multi-Statement Planner

Backup and Recovery (PITR) Requirement

- Transaction status should be consistent
  - Each transaction must be either:
    - Committed in all the involved node
    - Running or aborted in all the involved node
- Write such timing in WALs of all the coordinators and datanodes.
- Application can provide such timing as “BARRIER”
  - CREATE BARRIER *barrier_id*
    - Wait partially-committed-transactions completes commit,
    - Block other transaction’s commit,
    - Write BARRIER record to WALs of all the coordinators/datanodes.
  - When running PITR, specify barrier_id in recovery.conf

Global IT Innovator
NTT DATA GROUP

Demonstration

Further Development Topics/Schedule（1）

- Support more variety of statements:
  - HAVING, PREPARE, EXECUTE, CURSOR, TRIGGER
    - By the end of this year
  - SAVEPOINT
  - Multi-statement planner for WITH, WITH RECURSIVE, general functions, general subqueries, SELECT INTO, CREATE TABLE AS
    - By the end of this year

Further Development Topics/Schedule (2)

- Datanode high-availability
  - Backup with synchronous streaming replication
    - Synchronous replication needed to maintain data integrity among datanodes.
- Cluster operation
  - Online server addition/removal
- Challenging
  - Global constraint
    - Unique/Reference integrity among partition,
    - Exclusion constraint among partition
  - LOB
- Others needs additional test
  - dblink
  - SQL/MED

Postgres-XC to PostgreSQL

- Snapshot cloning
  - Parallel pg_dump
  - Parallel query execution (local/cluster)
- SQL/MED extension
  - Column projection pushdown
  - Join pushdown
  - Function pushdown
- Federation
  - Materialization
  - Cross-node join
  - Cross-node aggregation

Many candidate features.
Need more members for quick actions.

New Developer Wanted

- Writing Code
  - New distributed/parallel query handling/optimization
  - HA capabilities
  - Utilities
    - Installation
    - Configuration
    - Operation
  - Bug fixes
  - Back port to PostgreSQL
- Build
  - Creating binaries/distribution packages
- Test
  - Performance evaluation with various benchmarks
  - Finding bugs
  - New feature proposals
- Pilot application
  - Practical applications

Project resources

- Development site
  - http://sourceforge.net/projects/postgres-xc/
  - http://sourceforge.net/apps/mediawiki/postgres-xc/
- Project home
  - http://postgres-xc.sourceforge.net/
- Mailing List
  - http://postgres-xc.sourceforge.net/mailinglist.html

Contact us

Global IT Innovator
NTT DATA GROUP

Thank you very much!

Koichi Suzuki

NTT DATA INTELLILINK Corporation

Pacific Marks Tsukishima,1-15-7,
Tsukishima, Chuo-ku,
Tokyo 104-0052, Japan

TEL : +81 3 5843 6800
E-mail : koichi@intellilink.co.jp
koichi.szk@gmail.com
URL : http://www.intellilink.co.jp/ *only in Japanese

Global IT Innovator

NTT DATA GROUP